PySpark On AWS Glue

Creating Resources using supplied cloud formation template in aws

A **CloudFormation template** is a file (YAML or JSON) that describes:

- **What AWS resources to create**
- **How they are configured**
- **How they relate to each other**

When you choose **Create resources using supplied CloudFormation template**, AWS uses that template to create all required resources **in one operation**.

Go to AWS console and type cloud formation

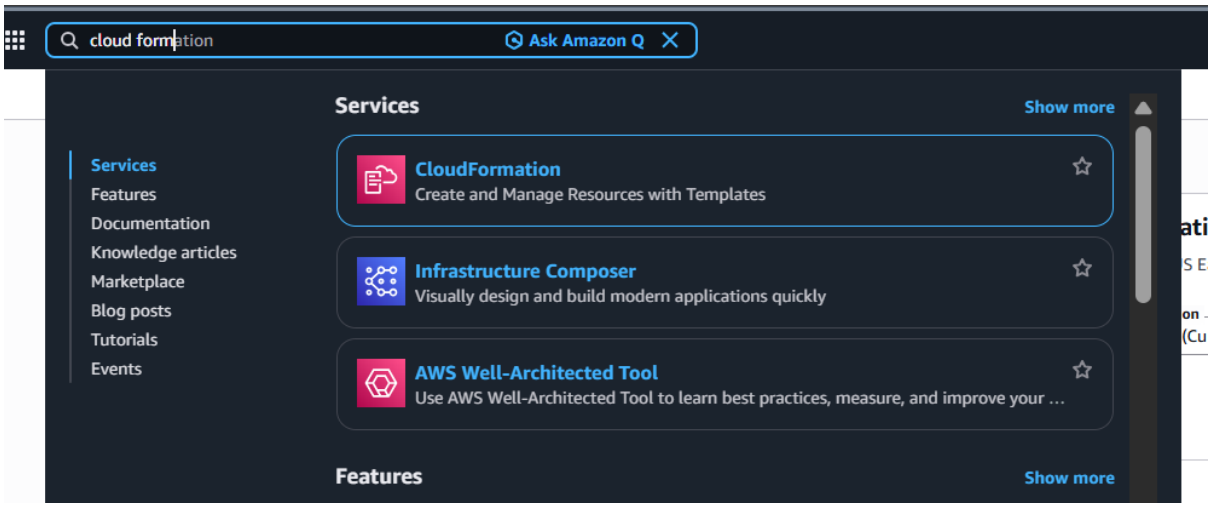

U must be a user that has permissions to run cloud formation,quick glue,resources create ion poilicies and create s3 buckets etc

I am using my IAM user account where I have my admin access

Now the cloudformation interface looks like the below, click on create stack

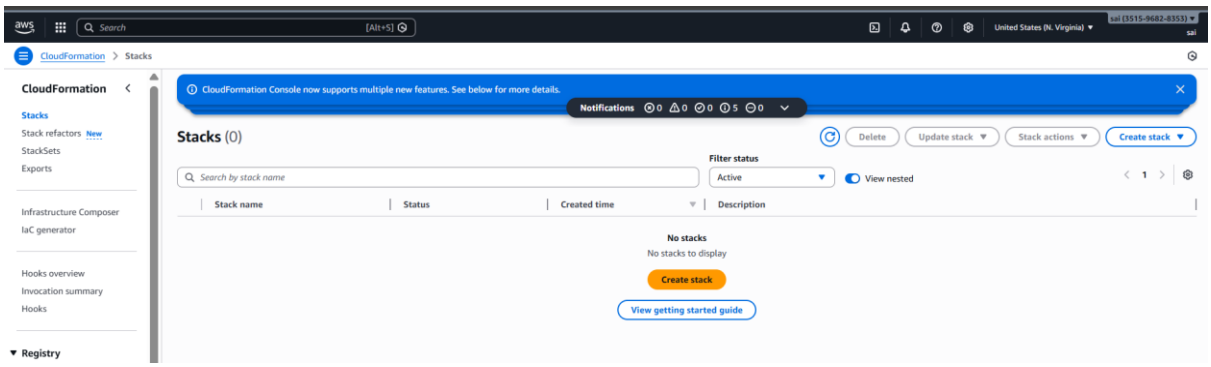

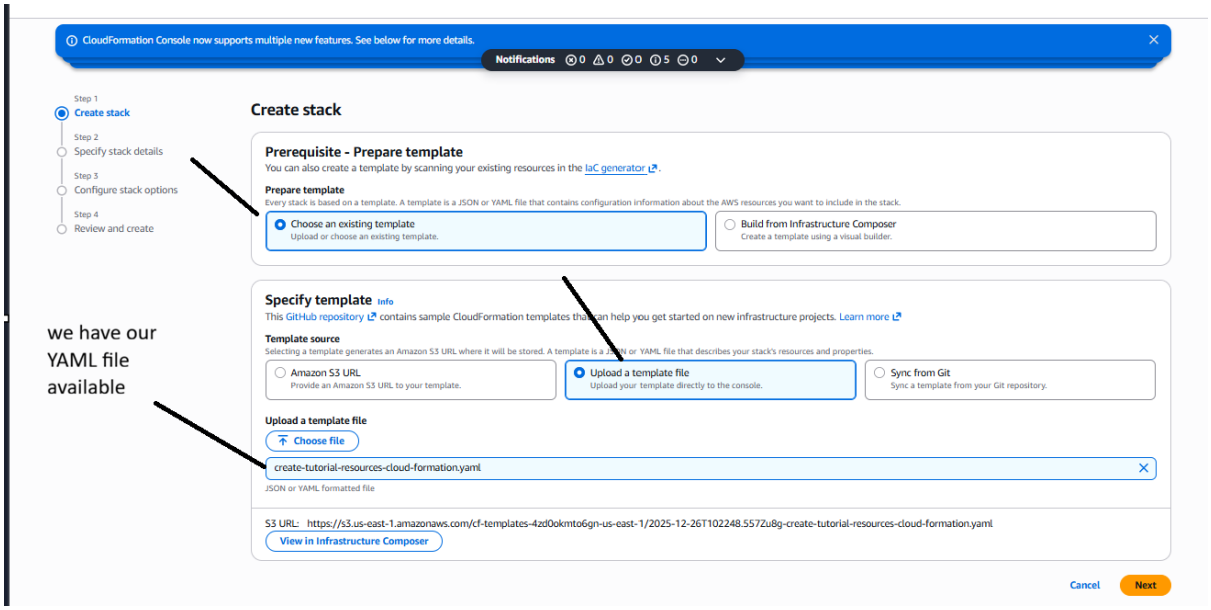

Click on next

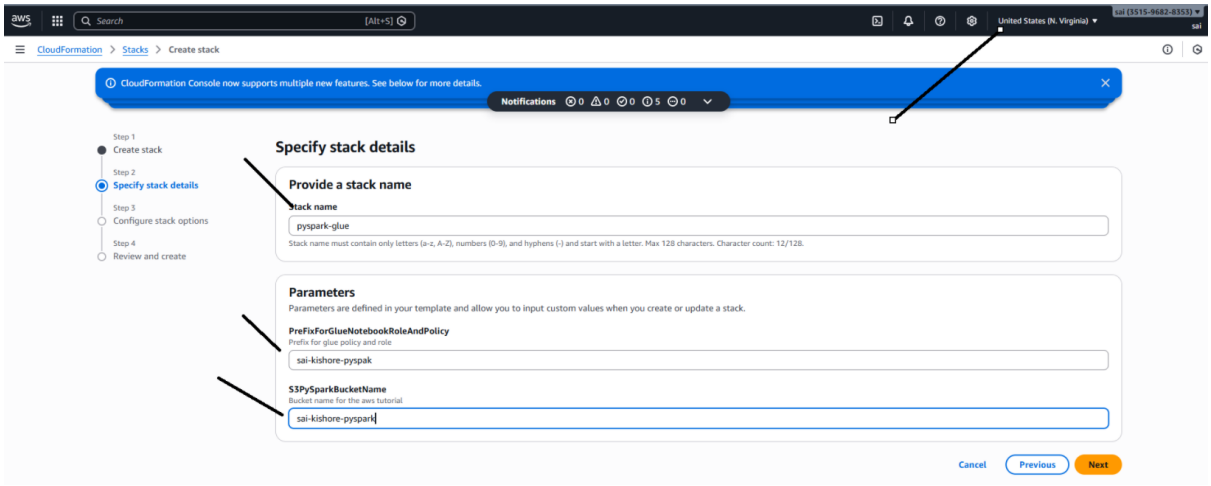

Click next

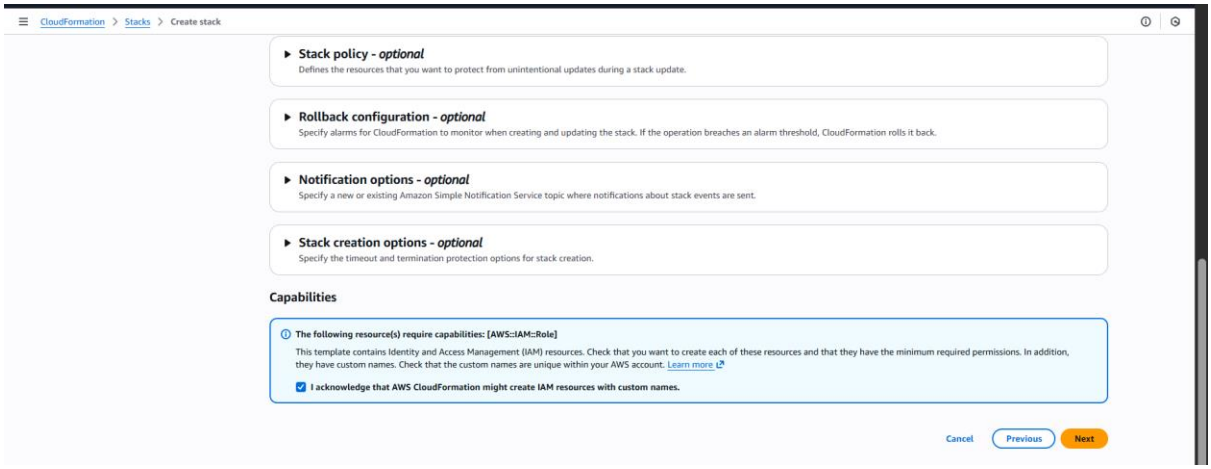

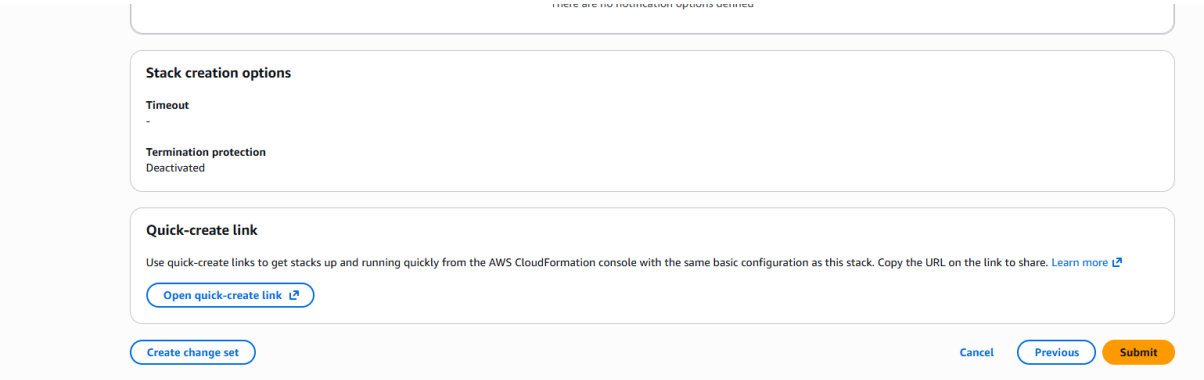

Click submit

The interface looks as follows

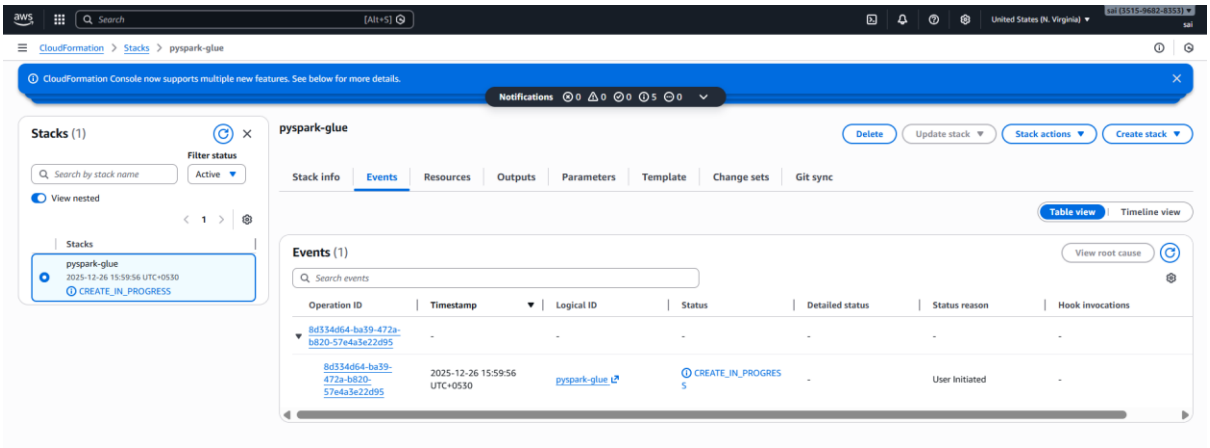

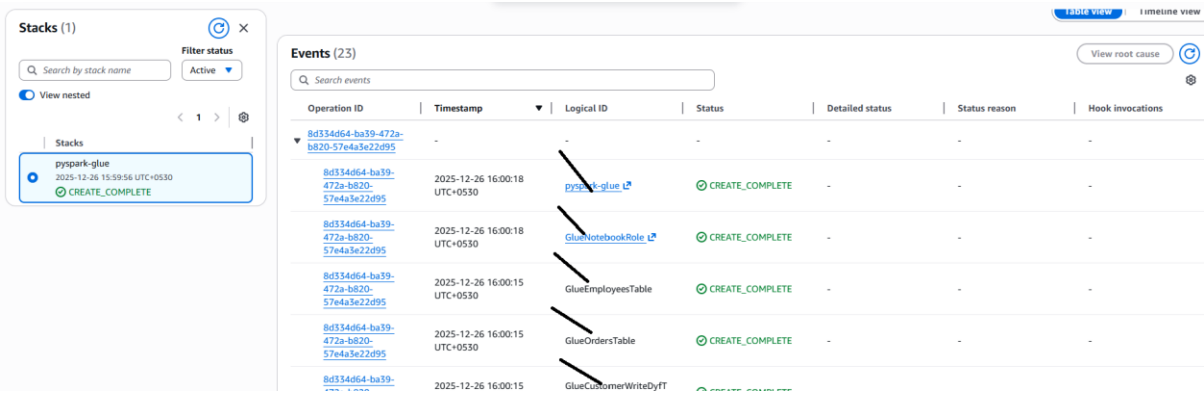

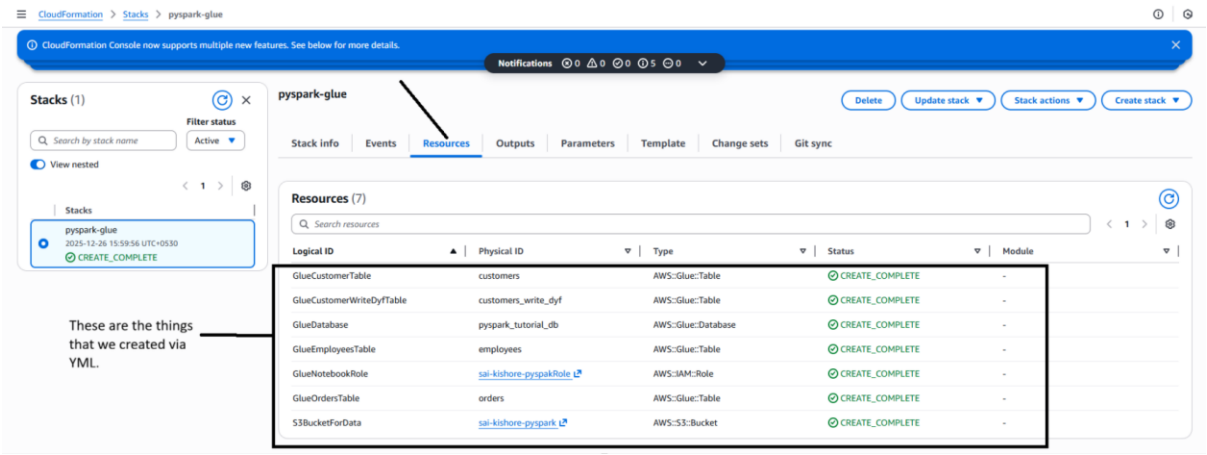

Now when u go to s3

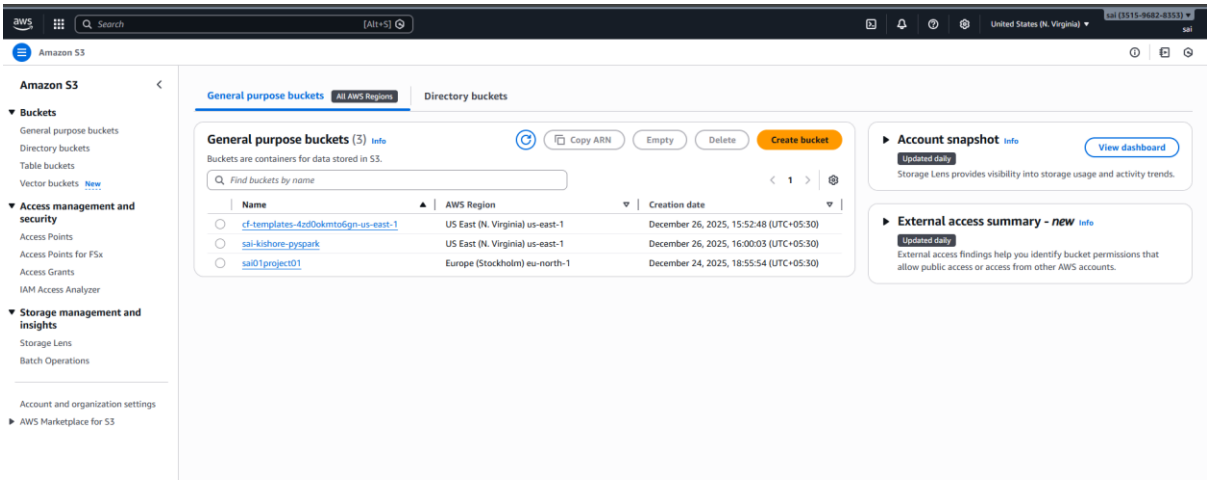

Uploaded folders with each file in it in csv format in sai-kishore-pyspark

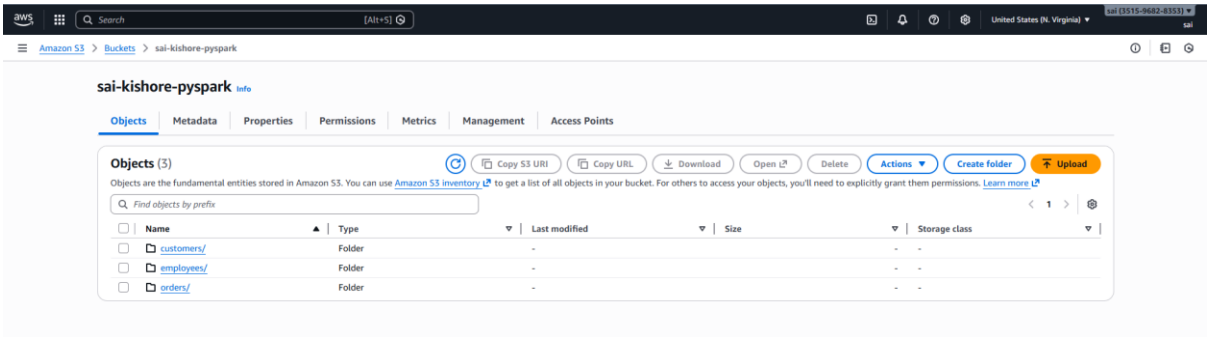

Now lets create a notebook server that we can use the glue interactive session.

Lets go to glue

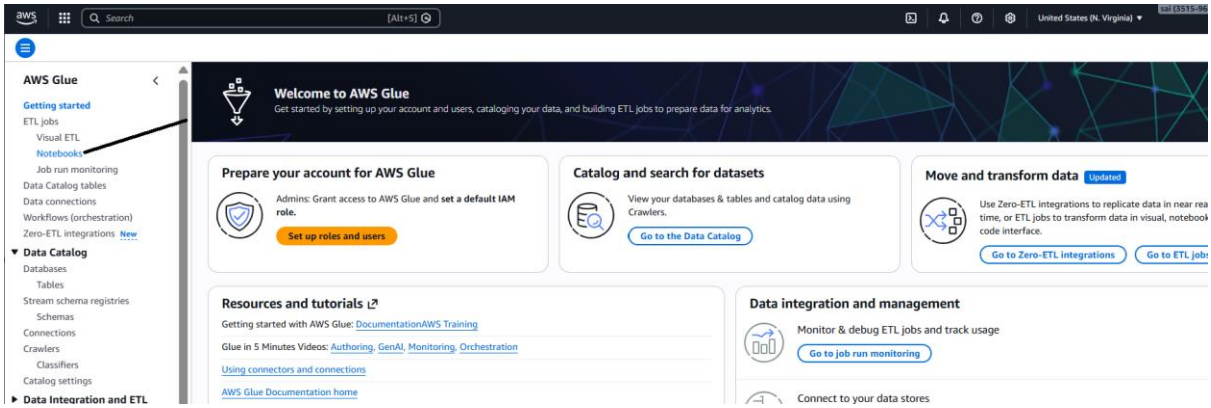

Click on notebook in below img

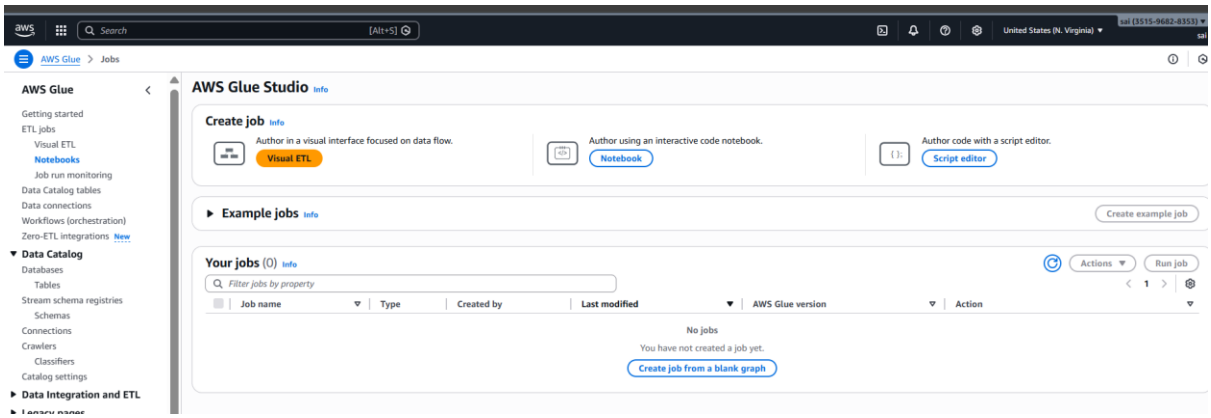

Click create notebook

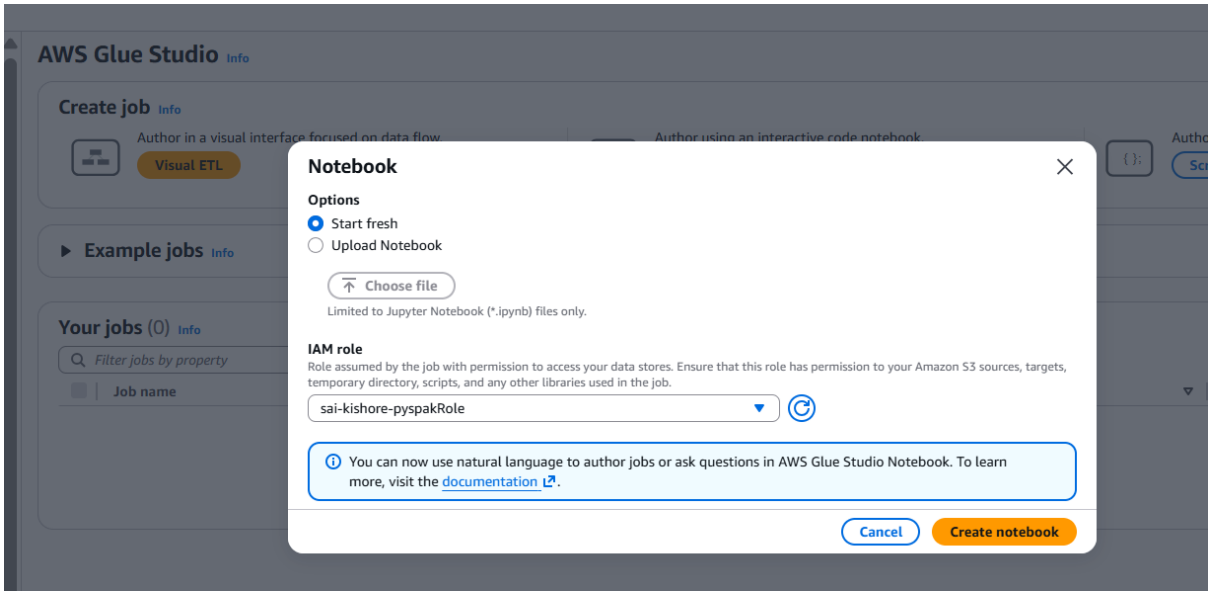

U see the interface like this

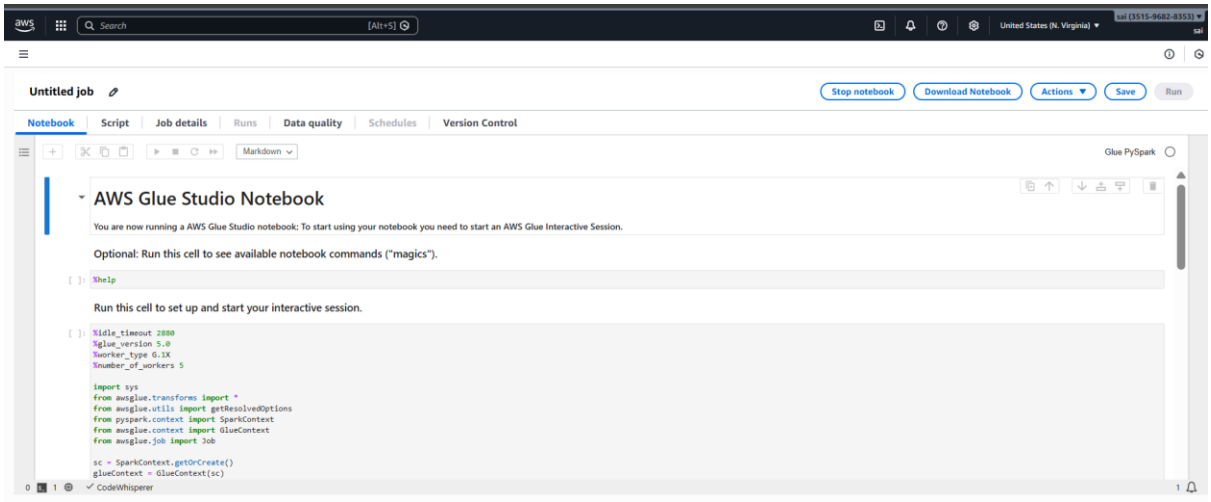


```
sc = SparkContext.getOrCreate()
glueContext = GlueContext(sc)
spark = glueContext.spark_session
job = Job(glueContext)
```

Example: Create a DynamicFrame from a table in the AWS Glue Data Catalog and display its schema

```
[ ]: dyf = glueContext.create_dynamic_frame.from_catalog(database='database_name', table_name='table_name')
dyf.printSchema()
```

Example: Convert the DynamicFrame to a Spark DataFrame and display a sample of the data

```
[ ]: df = dyf.toDF()
df.show()
```

Example: Visualize data with matplotlib

```
[ ]: import matplotlib.pyplot as plt

# Set X-axis and Y-axis values
x = [5, 2, 8, 4, 9]
y = [10, 4, 8, 5, 2]

# Create a bar chart
plt.bar(x, y)

# Show the plot
%matplot plt
```

Example: Write the data in the DynamicFrame to a location in Amazon S3 and a table for it in the AWS Glue Data Catalog ¶

```
[ ]: s3output = glueContext.getSink(
  path="s3://bucket_name/folder_name",
  connection_type="s3",
  updateBehavior="UPDATE_IN_DATABASE",
  partitionKeys=[],
  compression="snappy",
  enableUpdateCatalog=True,
  transformation_ctx="s3output",
)
s3output.setCatalogInfo(
  catalogDatabase="demo", catalogTableName="populations"
```

```
  transformation_ctx="s3output",
)
s3output.setCatalogInfo(
  catalogDatabase="demo", catalogTableName="populations"
)
s3output.setFormat("glueparquet")
s3output.writeFrame(DyF)
```

The above block of code appears by default

When I run this cell, spark session is created

Run this cell to set up and start your interactive session.

```
[1]: %idle_timeout 2880
%glue_version 5.0
%worker_type G.1X
%number_of_workers 5

import sys
from awsglue.transforms import *
from awsglue.utils import getResolvedOptions
from pyspark.context import SparkContext
from awsglue.context import GlueContext
from awsglue.job import Job

sc = SparkContext.getOrCreate()
glueContext = GlueContext(sc)
spark = glueContext.spark_session
job = Job(glueContext)
```

```
Welcome to the Glue Interactive Sessions Kernel
For more information on available magic commands, please type %help in any new cell.

Please view our Getting Started page to access the most up-to-date information on the Interactive Sessions kernel: https://docs.aws.amazon.com/glue/latest/dg/interactive-sessions.html
Installed kernel version: 1.0.7
Current idle_timeout is None minutes.
idle_timeout has been set to 2880 minutes.
Setting Glue version to: 5.0
Previous worker type: None
Setting new worker type to: G.1X
Previous number of workers: None
Setting new number of workers to: 5
Trying to create a Glue session for the kernel.
Session Type: glueet1
Worker Type: G.1X
Number of Workers: 5
Idle Timeout: 2880
Session ID: 3049a25d-6f5c-41ee-aed4-1e98e2258830
Applying the following default arguments:
--glue_kernel_version 1.0.7
--enable-glue-datacatalog true
Waiting for session 3049a25d-6f5c-41ee-aed4-1e98e2258830 to get into ready status...
Session 3049a25d-6f5c-41ee-aed4-1e98e2258830 has been created.
```

Example: Create a DynamicFrame from a table in the AWS Glue Data Catalog and display its schema

```
[ ]: dyf = glueContext.create_dynamic_frame.from_catalog(database='database_name', table_name='table_name')
dyf.printSchema()
```

```
[3]: # Read from the customers table in the glue data catalog using a dynamic frame
dynamicFrameCustomers = glueContext.create_dynamic_frame.from_catalog(
database = "pyspark_tutorial_db",
table_name = "customers"
)

# Show the top 10 rows from the dynamic dataframe
dynamicFrameCustomers.show(10)
```

```
{"customerid": 293, "firstname": "Catherine", "lastname": "Abel", "fullname": "Catherine Abel"}
{"customerid": 295, "firstname": "Kim", "lastname": "Abercrombie", "fullname": "Kim Abercrombie"}
{"customerid": 297, "firstname": "Humberto", "lastname": "Acevedo", "fullname": "Humberto Acevedo"}
{"customerid": 291, "firstname": "Gustavo", "lastname": "Achong", "fullname": "Gustavo Achong"}
{"customerid": 299, "firstname": "Pilar", "lastname": "Ackerman", "fullname": "Pilar Ackerman"}
{"customerid": 305, "firstname": "Carla", "lastname": "Adams", "fullname": "Carla Adams"}
{"customerid": 301, "firstname": "Frances", "lastname": "Adams", "fullname": "Frances Adams"}
{"customerid": 307, "firstname": "Jay", "lastname": "Adams", "fullname": "Jay Adams"}
{"customerid": 309, "firstname": "Ronald", "lastname": "Adina", "fullname": "Ronald Adina"}
```

1 Initialized (additional servers needed) Glue PySpark Idle CodeWhisperer

GLUE INTERACTIVE SESSIONS

- A programmatic and visual interface for building and testing extract, transform, and load (ETL) scripts for data preparation.
- Interactive sessions run Apache Spark analytics applications and provide on-demand access to a remote Spark runtime environment.
- AWS Glue transparently manages serverless Spark for these interactive

FUNDAMENTALS OF SPARK FOR GLUE

Apache Spark is an open-source in memory distributed processing system used for big data workloads

Apache Spark is an open-source **in memory** distributed processing system used for big data workloads

Apache Spark is an open-source in memory **distributed processing** system used for big data workloads

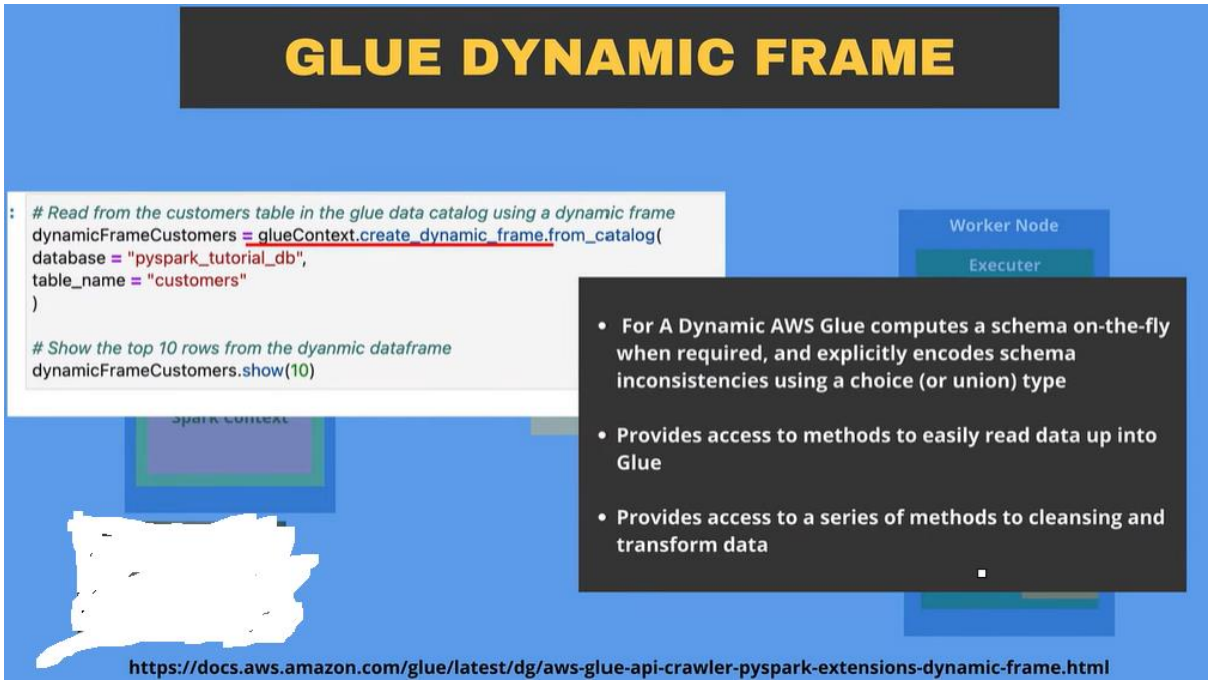


```
# Show the top 10 rows from the dyanmic dataframe
dynamicFrameCustomers.show(10)
```


```
{"customerid": 293, "firstname": "Catherine", "lastname": "Abel", "fullname": "Catherine Abel"}
{"customerid": 295, "firstname": "Kim", "lastname": "Abercrombie", "fullname": "Kim Abercrombie"}
{"customerid": 297, "firstname": "Humberto", "lastname": "Acevedo", "fullname": "Humberto Acevedo"}
{"customerid": 291, "firstname": "Gustavo", "lastname": "Achong", "fullname": "Gustavo Achong"}
{"customerid": 299, "firstname": "Pilar", "lastname": "Ackerman", "fullname": "Pilar Ackerman"}
{"customerid": 305, "firstname": "Carla", "lastname": "Adams", "fullname": "Carla Adams"}
{"customerid": 301, "firstname": "Frances", "lastname": "Adams", "fullname": "Frances Adams"}
{"customerid": 307, "firstname": "Jay", "lastname": "Adams", "fullname": "Jay Adams"}
{"customerid": 309, "firstname": "Ronald", "lastname": "Adina", "fullname": "Ronald Adina"}
{"customerid": 311, "firstname": "Samuel", "lastname": "Agcaoili", "fullname": "Samuel Agcaoili"}
{"customerid": 313, "firstname": "James", "lastname": "Aguilar", "fullname": "James Aguilar"}
{"customerid": 315, "firstname": "Robert", "lastname": "Ahlering", "fullname": "Robert Ahlering"}
{"customerid": 319, "firstname": "Kim", "lastname": "Akers", "fullname": "Kim Akers"}
{"customerid": 441, "firstname": "Stanley", "lastname": "Alan", "fullname": "Stanley Alan"}
{"customerid": 323, "firstname": "Amy", "lastname": "Alberts", "fullname": "Amy Alberts"}
{"customerid": 325, "firstname": "Anna", "lastname": "Albright", "fullname": "Anna Albright"}
{"customerid": 327, "firstname": "Milton", "lastname": "Albury", "fullname": "Milton Albury"}
{"customerid": 329, "firstname": "Paul", "lastname": "Alcorn", "fullname": "Paul Alcorn"}
{"customerid": 331, "firstname": "Gregory", "lastname": "Alderson", "fullname": "Gregory Alderson"}
{"customerid": 333, "firstname": "J. Phillip", "lastname": "Alexander", "fullname": "J. Phillip Alexander"}
```


We call it dynamic frame when we have many records together.

Aws gives us numerous methods to use with a dynamic frame

Dynamic frame has different methods to read the data like

RDD

JDBC

S3

Glue Data Catalog – from_catalog is used in the code

Here is the documentation link for Dynamicframe class:

https://docs.aws.amazon.com/glue/latest/dg/aws-glue-api-crawler-pyspark-extensions-dynamic-frame.html

U will also have practical examples in it, so go through documentation

Let's start our code

Print schema: S is capital in code

```
{ customerid : 301,  firstname :  Frances ,  lastname :  Adams ,  fullname :  Frances Adams }
{"customerid": 307, "firstname": "Jay", "lastname": "Adams", "fullname": "Jay Adams"}
{"customerid": 309, "firstname": "Ronald", "lastname": "Adina", "fullname": "Ronald Adina"}
{"customerid": 311, "firstname": "Samuel", "lastname": "Agcaoili", "fullname": "Samuel Agcaoili"}
```

```
[3]: #print schema
dynamicFrameCustomers.printSchema()
```

```
root
|-- customerid: long
|-- firstname: string
|-- lastname: string
|-- fullname: string
```

```
|-- lastname: string
|-- fullname: string
```

```
[4]: #count
dynamicFrameCustomers.count()
```

```
635
```

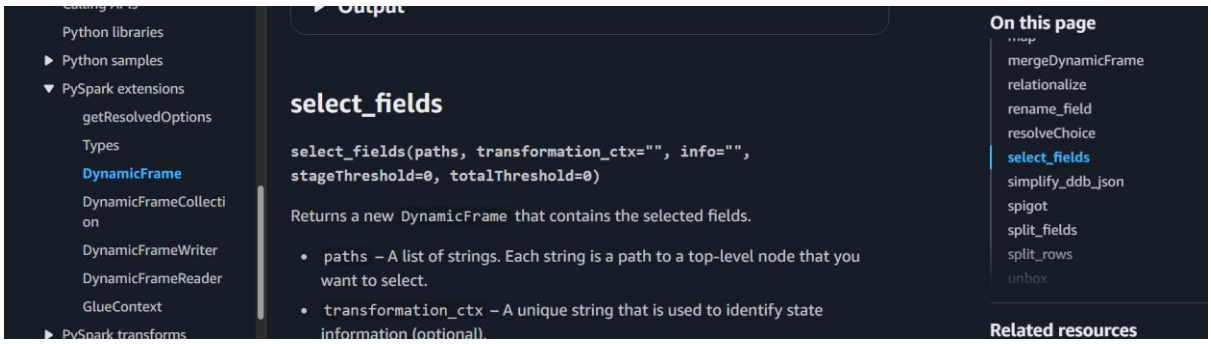


```
[8]: #selecting customerid and full name
dyfselectcustomers = dynamicFrameCustomers.select_fields(["customerid","fullname"])

#show the result
dyfselectcustomers.show(10)
```

```
{"customerid": 293, "fullname": "Catherine Abel"}
{"customerid": 295, "fullname": "Kim Abercrombie"}
{"customerid": 297, "fullname": "Humberto Acevedo"}
{"customerid": 291, "fullname": "Gustavo Achong"}
{"customerid": 299, "fullname": "Pilar Ackerman"}
{"customerid": 305, "fullname": "Carla Adams"}
{"customerid": 301, "fullname": "Frances Adams"}
{"customerid": 307, "fullname": "Jay Adams"}
{"customerid": 309, "fullname": "Ronald Adina"}
{"customerid": 311, "fullname": "Samuel Agcaoili"}
```

Note: always follow the documentation if u need any help

[11]:
```
dynamicFrameCustomers.select_fields(["customerid","fullname"]).show(10)
```

```
{"customerid": 293, "fullname": "Catherine Abel"}
{"customerid": 295, "fullname": "Kim Abercrombie"}
{"customerid": 297, "fullname": "Humberto Acevedo"}
{"customerid": 291, "fullname": "Gustavo Achong"}
{"customerid": 299, "fullname": "Pilar Ackerman"}
{"customerid": 305, "fullname": "Carla Adams"}
{"customerid": 301, "fullname": "Frances Adams"}
{"customerid": 307, "fullname": "Jay Adams"}
{"customerid": 309, "fullname": "Ronald Adina"}
{"customerid": 311, "fullname": "Samuel Agcaoili"}
```

[12]:
```
#Drop Fields of Dynamic Frame
dyfCustomerDropFields = dynamicFrameCustomers.drop_fields(["firstname","lastname"])

# Show Top 10 rows of dyfCustomerDropFields Dynamic Frame
dyfCustomerDropFields.show(10)
```

```
{"customerid": 293, "fullname": "Catherine Abel"}
{"customerid": 295, "fullname": "Kim Abercrombie"}
{"customerid": 297, "fullname": "Humberto Acevedo"}
{"customerid": 291, "fullname": "Gustavo Achong"}
{"customerid": 299, "fullname": "Pilar Ackerman"}
{"customerid": 305, "fullname": "Carla Adams"}
{"customerid": 301, "fullname": "Frances Adams"}
{"customerid": 307, "fullname": "Jay Adams"}
{"customerid": 309, "fullname": "Ronald Adina"}
{"customerid": 311, "fullname": "Samuel Agcaoili"}
```

[13]:
```
# Mapping array for column rename fullname -> name
mapping=[("customerid", "long", "customerid","long"),("fullname", "string", "name", "string")]

# Apply the mapping to rename fullname -> name
dfyMapping = ApplyMapping.apply(
                                frame = dyfCustomerDropFields,
                                mappings = mapping,
                                transformation_ctx = "applymapping1"
                                )

# show the new dynamic frame with name column
dfyMapping.show(10)
```

```
{"customerid": 293, "name": "Catherine Abel"}
{"customerid": 295, "name": "Kim Abercrombie"}
{"customerid": 297, "name": "Humberto Acevedo"}
{"customerid": 291, "name": "Gustavo Achong"}
{"customerid": 299, "name": "Pilar Ackerman"}
{"customerid": 305, "name": "Carla Adams"}
{"customerid": 301, "name": "Frances Adams"}
{"customerid": 307, "name": "Jay Adams"}
{"customerid": 309, "name": "Ronald Adina"}
{"customerid": 311, "name": "Samuel Agcaoili"}
```

[ ]:

[14]:
```
# Filter dynamicFrameCustomers for customers who have the last name Adams
dyfFilter=  Filter.apply(frame = dynamicFrameCustomers,
                                        f = lambda x: x["lastname"] in "Adams"
                                    )

# Show the top 10 customers  from the filtered Dynamic frame
dyfFilter.show(10)
```

```
{"lastname": "Adams", "firstname": "Carla", "customerid": 305, "fullname": "Carla Adams"}
{"lastname": "Adams", "firstname": "Frances", "customerid": 301, "fullname": "Frances Adams"}
{"lastname": "Adams", "firstname": "Jay", "customerid": 307, "fullname": "Jay Adams"}
```

```
[15]: # Read from the customers table in the glue data catalog using a dynamic frame
dynamicFrameOrders = glueContext.create_dynamic_frame.from_catalog(
database = "pyspark_tutorial_db",
table_name = "orders"
)

# show top 10 rows of orders table
dynamicFrameOrders.show(10)
```


{"salesorderid": 43659, "salesorderdetailid": 1, "orderdate": "5/31/2011", "duedate": "6/12/2011", "shipdate": "6/7/2011", "employeeid": 279, "customerid": 1045, "subtotal": 20565.6206, "taxamt": 1971.5149, "freight": 616.0984, "totaldue": 23153.2339, "productid": 776, "orderqty": 1, "unitprice": 2024.9940, "unitpricediscount": 0.0000, "linetotal": 2024.9940}
{"salesorderid": 43659, "salesorderdetailid": 2, "orderdate": "5/31/2011", "duedate": "6/12/2011", "shipdate": "6/7/2011", "employeeid": 279, "customerid": 1045, "subtotal": 20565.6206, "taxamt": 1971.5149, "freight": 616.0984, "totaldue": 23153.2339, "productid": 777, "orderqty": 3, "unitprice": 2024.9940, "unitpricediscount": 0.0000, "linetotal": 6074.9820}
{"salesorderid": 43659, "salesorderdetailid": 3, "orderdate": "5/31/2011", "duedate": "6/12/2011", "shipdate": "6/7/2011", "employeeid": 279, "customerid": 1045, "subtotal": 20565.6206, "taxamt": 1971.5149, "freight": 616.0984, "totaldue": 23153.2339, "productid": 778, "orderqty": 1, "unitprice": 2024.9940, "unitpricediscount": 0.0000, "linetotal": 2024.9940}
{"salesorderid": 43659, "salesorderdetailid": 4, "orderdate": "5/31/2011", "duedate": "6/12/2011", "shipdate": "6/7/2011", "employeeid": 279, "customerid": 1045, "subtotal": 20565.6206, "taxamt": 1971.5149, "freight": 616.0984, "totaldue": 23153.2339, "productid": 771, "orderqty": 1, "unitprice": 2039.9940, "unitpricediscount": 0.0000, "linetotal": 2039.9940}
{"salesorderid": 43659, "salesorderdetailid": 5, "orderdate": "5/31/2011", "duedate": "6/12/2011", "shipdate": "6/7/2011", "employeeid": 279, "customerid": 1045, "subtotal": 20565.6206, "taxamt": 1971.5149, "freight": 616.0984, "totaldue": 23153.2339, "productid": 772, "orderqty": 1, "unitprice": 2039.9940, "unitpricediscount": 0.0000, "linetotal": 2039.9940}
{"salesorderid": 43659, "salesorderdetailid": 6, "orderdate": "5/31/2011", "duedate": "6/12/2011", "shipdate": "6/7/2011", "employeeid": 279, "customerid": 1045, "subtotal": 20565.6206, "taxamt": 1971.5149, "freight": 616.0984, "totaldue": 23153.2339, "productid": 773, "orderqty": 2, "unitprice": 2039.9940, "unitpricediscount": 0.0000, "linetotal": 4079.9880}
{"salesorderid": 43659, "salesorderdetailid": 7, "orderdate": "5/31/2011", "duedate": "6/12/2011", "shipdate": "6/7/2011", "employeeid": 279, "customerid": 1045, "subtotal": 20565.6206, "taxamt": 1971.5149, "freight": 616.0984, "totaldu

Note: u see that the interface in athena of ur tables in s3, means the data got crawled by yaml file and data crawling is done in prev AWS project

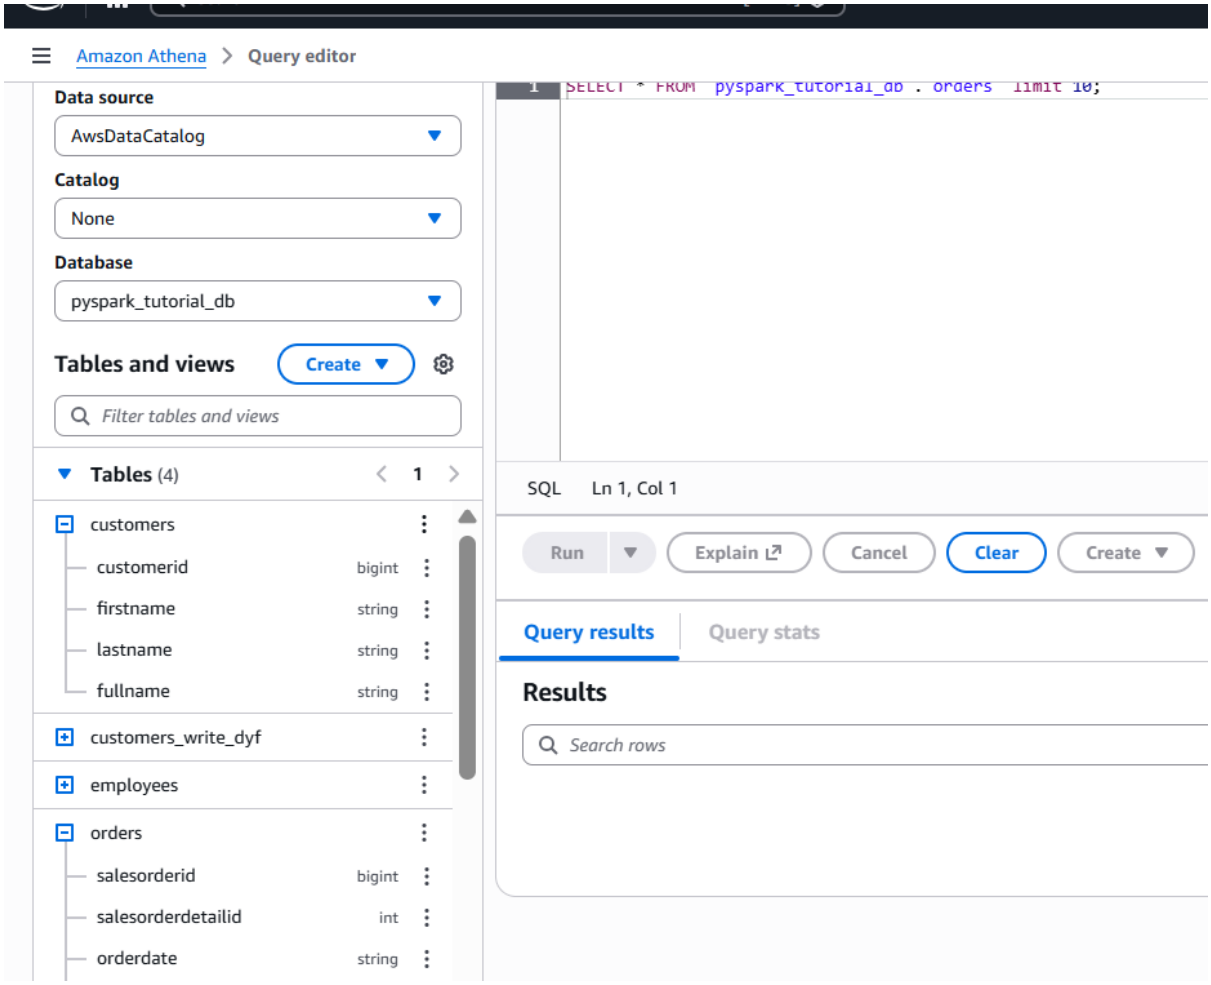


```
[16]: # Join two dynamic frames on an equality join
dyfjoin = dynamicFrameCustomers.join(["customerid"],["customerid"],dynamicFrameOrders)

# show top 10 rows for the joined dynamic
dyfjoin.show(10)
```


{"freight": 181.0019, "subtotal": 6035.8246, "salesorderdetailid": 1628, "productid": 754, "linetotal": 874.7940, "employeeid": 277, ".customerid": 671, "taxamt": 579.2061, "salesorderid": 44097, "duedate": "8/13/2011", "orderqty": 1, "shipdate": "8/8/2011", "lastname": "Chapla", "firstname": "Lee", "totaldue": 6796.0326, "unitprice": 874.7940, "orderdate": "8/1/2011", "unitpricediscount": 0.0000, "customerid": 671, "fullname": "Lee Chapla"}
{"freight": 181.0019, "subtotal": 6035.8246, "salesorderdetailid": 1629, "productid": 760, "linetotal": 419.4589, "employeeid": 277, ".customerid": 671, "taxamt": 579.2061, "salesorderid": 44097, "duedate": "8/13/2011", "orderqty": 1, "shipdate": "8/8/2011", "lastname": "Chapla", "firstname": "Lee", "totaldue": 6796.0326, "unitprice": 419.4589, "orderdate": "8/1/2011", "unitpricediscount": 0.0000, "customerid": 671, "fullname": "Lee Chapla"}
{"freight": 181.0019, "subtotal": 6035.8246, "salesorderdetailid": 1630, "productid": 762, "linetotal": 838.9178, "employeeid": 277, ".customerid": 671, "taxamt": 579.2061, "salesorderid": 44097, "duedate": "8/13/2011", "orderqty": 2, "shipdate": "8/8/2011", "lastname": "Chapla", "firstname": "Lee", "totaldue": 6796.0326, "unitprice": 419.4589, "orderdate": "8/1/2011", "unitpricediscount": 0.0000, "customerid": 671, "fullname": "Lee Chapla"}
{"freight": 181.0019, "subtotal": 6035.8246, "salesorderdetailid": 1631, "productid": 708, "linetotal": 80.7460, "employeeid": 277, ".customerid": 671, "taxamt": 579.2061, "salesorderid": 44097, "duedate":

Creating a folder name “write_down_dyf_to_s3”

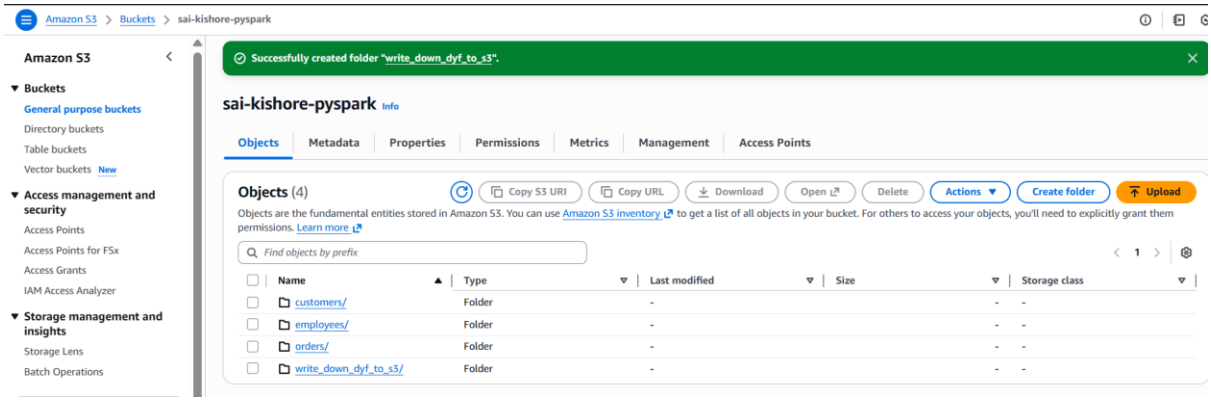


```
671, "fullname": "Lee Chapla"}
```

```
[17]: # write down the data in a Dynamic Frame to S3 location.
glueContext.write_dynamic_frame.from_options(
                        frame = dynamicFrameCustomers,
                        connection_type="s3",
                        connection_options = {"path": "s3://sai-kishore-pyspark/write_down_dyf_to_s3/"},
                        format = "csv",
                        format_options={
                            "separator": ","
                            },
                        transformation_ctx = "datasink2")
```

```
<awsglue.dynamicframe.DynamicFrame object at 0x7f277b428550>
```

```
[ ]:
```

Now check ur s3 bucket

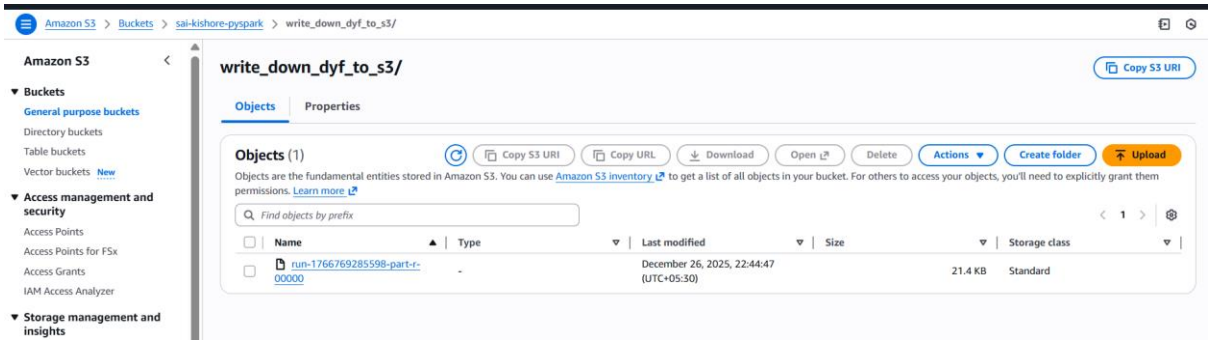

Download the file and check the data

```
C: > Users > GudiboinaSaiKishore > Downloads > run-1766769285598-part-r-00000
customerid,firstname,lastname,fullname
293,Catherine,Abel,"Catherine Abel"
295,Kim,Abercrombie,"Kim Abercrombie"
297,Humberto,Acevedo,"Humberto Acevedo"
291,Gustavo,Achong,"Gustavo Achong"
299,Pilar,Ackerman,"Pilar Ackerman"
305,Carla,Adams,"Carla Adams"
301,Frances,Adams,"Frances Adams"
307,Jay,Adams,"Jay Adams"
309,Ronald,Adina,"Ronald Adina"
311,Samuel,Agcaoili,"Samuel Agcaoili"
313,James,Aguilar,"James Aguilar"
315,Robert,Ahlering,"Robert Ahlering"
319,Kim,Akers,"Kim Akers"
441,Stanley,Alan,"Stanley Alan"
323,Amy,Alberts,"Amy Alberts"
325,Anna,Albright,"Anna Albright"
327,Milton,Albury,"Milton Albury"
329,Paul,Alcorn,"Paul Alcorn"
331,Gregory,Alderson,"Gregory Alderson"
333,"J. Phillip",Alexander,"J. Phillip Alexander"
1149,Mary,Alexander,"Mary Alexander"
335,Michelle,Alexander,"Michelle Alexander"
341,Marvin,Allen,"Marvin Allen"
343,Michael,Allen,"Michael Allen"
339,Phyllis,Allen,"Phyllis Allen"
345,Cecil,Allison,"Cecil Allison"
347,Oscar,Alpuerto,"Oscar Alpuerto"
```

```
<awsglue.dynamicframe.DynamicFrame object at 0x7f277b428550>

[20]: # write data from the dynamicFrameCustomers to customers_write_dyf table using the meta data stored in the glue data catalog
glueContext.write_dynamic_frame.from_catalog(
    frame=dynamicFrameCustomers,
    database = "pyspark_tutorial_db",
    table_name = "customers_write_dyf"
)

<awsglue.dynamicframe.DynamicFrame object at 0x7f277b429810>

[ ]:
```

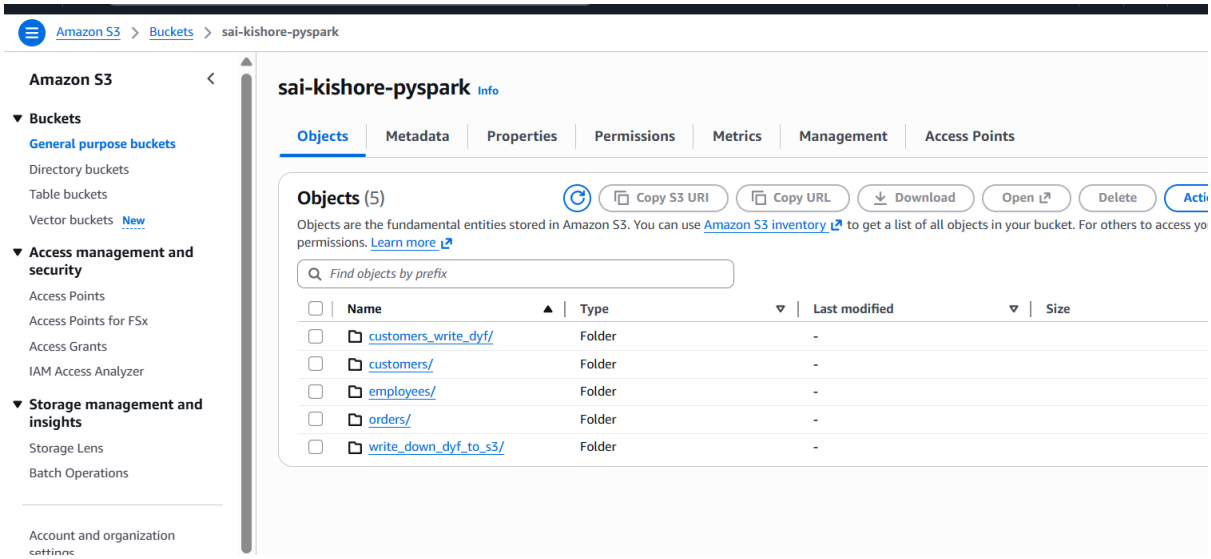

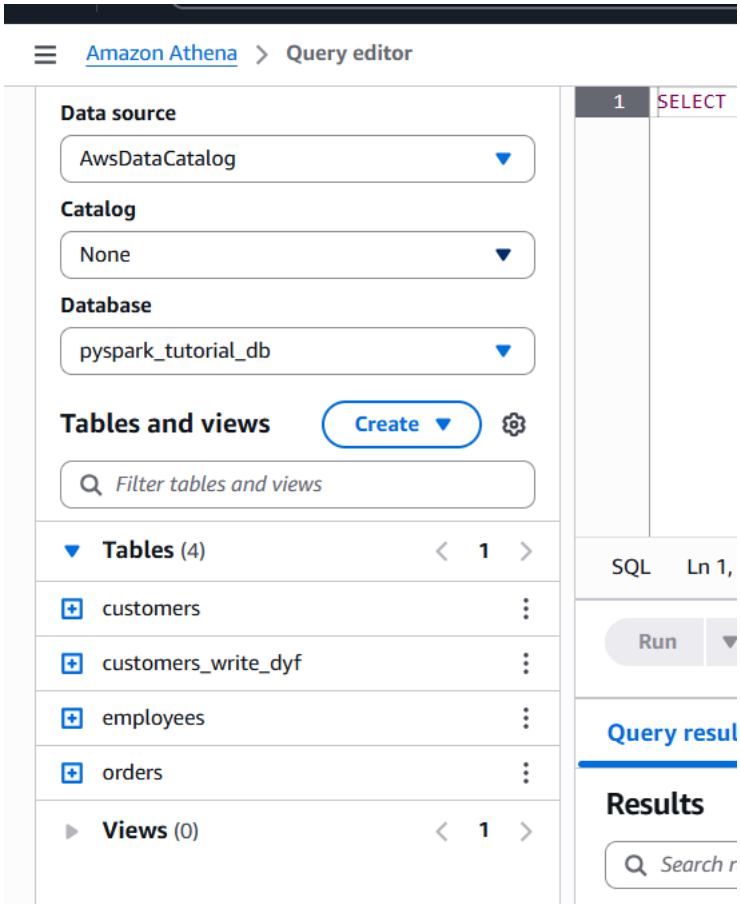

Note: write_down_dyf_to_s3 didn’t come to athena

[21]:
```
# Dynamic Frame to Spark DataFrame
sparkDf = dynamicFrameCustomers.toDF()

#show spark DF
sparkDf.show()
```

```
+----------+----------+-----------+--------------------+
|customerid| firstname|   lastname|            fullname|
+----------+----------+-----------+--------------------+
|       293| Catherine|       Abel|      Catherine Abel|
|       295|       Kim|Abercrombie|     Kim Abercrombie|
|       297|  Humberto|    Acevedo|    Humberto Acevedo|
|       291|   Gustavo|     Achong|      Gustavo Achong|
|       299|     Pilar|   Ackerman|      Pilar Ackerman|
|       305|     Carla|      Adams|         Carla Adams|
|       301|   Frances|      Adams|       Frances Adams|
|       307|       Jay|      Adams|           Jay Adams|
|       309|    Ronald|      Adina|        Ronald Adina|
|       311|    Samuel|   Agcaoili|     Samuel Agcaoili|
```

[22]:
```
# Select columns from spark dataframe
dfSelect = sparkDf.select("customerid","fullname")

# show selected
dfSelect.show()
```

```
+----------+--------------------+
|customerid|            fullname|
+----------+--------------------+
|       293|      Catherine Abel|
|       295|     Kim Abercrombie|
|       297|    Humberto Acevedo|
|       291|      Gustavo Achong|
|       299|      Pilar Ackerman|
|       305|         Carla Adams|
```

[23]:
```
#import lit from sql functions
from  pyspark.sql.functions import lit

# Add new column to spark dataframe
dfNewColumn = sparkDf.withColumn("date", lit("2022-07-24"))

# show df with new column
dfNewColumn.show()
```

```
+----------+----------+-----------+--------------------+----------+
|customerid| firstname|   lastname|            fullname|      date|
+----------+----------+-----------+--------------------+----------+
|       293| Catherine|       Abel|      Catherine Abel|2022-07-24|
|       295|       Kim|Abercrombie|     Kim Abercrombie|2022-07-24|
|       297|  Humberto|    Acevedo|    Humberto Acevedo|2022-07-24|
|       291|   Gustavo|     Achong|      Gustavo Achong|2022-07-24|
|       299|     Pilar|   Ackerman|      Pilar Ackerman|2022-07-24|
|       305|     Carla|      Adams|         Carla Adams|2022-07-24|
|       301|   Frances|      Adams|       Frances Adams|2022-07-24|
|       307|       Jay|      Adams|          Jay Adams|2022-07-24|
|       309|    Ronald|      Adina|       Ronald Adina|2022-07-24|
|       311|    Samuel|   Agcaoili|     Samuel Agcaoili|2022-07-24|
|       313|     James|    Aguilar|       James Aguilar|2022-07-24|
```

[24]:
```
#import concat from functions
from  pyspark.sql.functions import concat

# create another full name column
dfNewFullName = sparkDf.withColumn("new_full_name",concat("firstname",concat(lit(' '),"lastname")))

#show full name column
dfNewFullName.show()
```

```
+----------+----------+-----------+--------------------+--------------------+
|customerid| firstname|   lastname|            fullname|       new_full_name|
+----------+----------+-----------+--------------------+--------------------+
|       293| Catherine|       Abel|      Catherine Abel|      Catherine Abel|
|       295|       Kim|Abercrombie|     Kim Abercrombie|     Kim Abercrombie|
|       297|  Humberto|    Acevedo|    Humberto Acevedo|    Humberto Acevedo|
|       291|   Gustavo|     Achong|      Gustavo Achong|      Gustavo Achong|
|       299|     Pilar|   Ackerman|      Pilar Ackerman|      Pilar Ackerman|
|       305|     Carla|      Adams|         Carla Adams|         Carla Adams|
|       301|   Frances|      Adams|       Frances Adams|       Frances Adams|
|       307|       Jay|      Adams|          Jay Adams|          Jay Adams|
```

[25]:
```
# Drop column from spark dataframe
dfDropCol = sparkDf.drop("firstname","lastname")

#show dropped column df
dfDropCol.show()
```

```
+----------+--------------------+
|customerid|            fullname|
+----------+--------------------+
|       293|      Catherine Abel|
|       295|     Kim Abercrombie|
|       297|    Humberto Acevedo|
|       291|      Gustavo Achong|
|       299|      Pilar Ackerman|
|       305|         Carla Adams|
|       301|       Frances Adams|
|       307|          Jay Adams|
```

[26]:
```
# Rename column in Spark dataframe
dfRenameCol = sparkDf.withColumnRenamed("fullname","full_name")

#show renamed column dataframe
dfRenameCol.show()
```

```
+----------+----------+-----------+--------------------+
|customerid| firstname|   lastname|           full_name|
+----------+----------+-----------+--------------------+
|       293| Catherine|       Abel|      Catherine Abel|
|       295|       Kim|Abercrombie|     Kim Abercrombie|
|       297|  Humberto|    Acevedo|    Humberto Acevedo|
|       291|   Gustavo|     Achong|      Gustavo Achong|
|       299|     Pilar|   Ackerman|      Pilar Ackerman|
|       305|     Carla|      Adams|         Carla Adams|
|       301|   Frances|      Adams|       Frances Adams|
```

[27]:
```
# Group by lastname then print counts of lastname and show
sparkDf.groupBy("lastname").count().show()
```

```
+--------+-----+
|lastname|count|
+--------+-----+
|  Achong|    1|
|  Bailey|    1|
|   Caron|    1|
|   Casts|    1|
|   Curry|    1|
| Desalvo|    1|
| Dockter|    1|
|    Dyck|    1|
|  Farino|    1|
| Fluegel|    1|
|   Ganio|    1|
```

only showing top 20 rows

[28]:
```
# Filter spark DataFrame for customers who have the last name Adams
sparkDf.filter(sparkDf["lastname"] == "Adams").show()
```

```
+----------+---------+--------+-------------+
|customerid|firstname|lastname|     fullname|
+----------+---------+--------+-------------+
|       305|    Carla|   Adams|  Carla Adams|
|       301|  Frances|   Adams|Frances Adams|
|       307|      Jay|   Adams|    Jay Adams|
+----------+---------+--------+-------------+
```

[ ]:

[29]:
```
# Where clause spark DataFrame for customers who have the last name Adams
sparkDf.where("lastname ='Adams'").show()
```

```
+----------+---------+--------+-------------+
|customerid|firstname|lastname|     fullname|
+----------+---------+--------+-------------+
|       305|    Carla|   Adams|  Carla Adams|
|       301|  Frances|   Adams|Frances Adams|
|       307|      Jay|   Adams|    Jay Adams|
+----------+---------+--------+-------------+
```

[ ]:

[30]:
```
# Read from the customers table in the glue data catalog using a dynamic frame and convert to spark dataframe
dfOrders = glueContext.create_dynamic_frame.from_catalog(
                                    database = "pyspark_tutorial_db",
                                    table_name = "orders"
                                ).toDF()
```

[31]:
```
dfOrders.show(10)
```

```
+----------+-----------------+---------+---------+---------+----------+----------+----------+---------+---------+----------+---------+-----------+-----------------+---------+
|salesorderid|salesorderdetailid|orderdate|  duedate|shipdate|employeeid|customerid|  subtotal|   taxamt| freight|  totaldue|productid|orderqty|unitprice|unitpricediscount|linetotal|
+----------+-----------------+---------+---------+---------+----------+----------+----------+---------+---------+----------+---------+-----------+-----------------+---------+
|       43659|                1|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      776|          1|        2024.9940|           0.0000|2024.9940|
|       43659|                2|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      777|          3|        2024.9940|           0.0000|6074.9820|
|       43659|                3|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      778|          1|        2024.9940|           0.0000|2024.9940|
|       43659|                4|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      771|          1|        2039.9940|           0.0000|2039.9940|
|       43659|                5|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      772|          1|        2039.9940|           0.0000|2039.9940|
|       43659|                6|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      773|          2|        2039.9940|           0.0000|4079.9880|
|       43659|                7|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      774|          1|        2039.9940|           0.0000|2039.9940|
|       43659|                8|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      714|          3|  28.8404|           0.0000|  86.5212|
|       43659|                9|5/31/2011|6/12/2011|6/7/2011|       279|      1045|20565.6206|1971.5149|616.0984|23153.2339|      716|          1|  28.8404|           0.0000|  28.8404|
```

[32]:
```
# Inner Join Customers Spark DF to Orders Spark DF
sparkDf.join(dfOrders,sparkDf.customerid ==  dfOrders.customerid,"inner").show(truncate=False)
```

```
+----------+---------+--------+--------------+-----------+-----------------+---------+---------+-------+---------+----------+---------+-----------+---------+---------+
|customerid|firstname|lastname|fullname      |salesorderid|salesorderdetailid|orderdate|duedate  |shipdate|employeeid|customerid|subtotal  |taxamt   |freight |totaldue  |productid|orderqty|unitprice|unitpriediscount|linetotal|
+----------+---------+--------+--------------+-----------+-----------------+---------+---------+-------+---------+----------+---------+-----------+---------+---------+
|517       |Richard  |Bready  |Richard Bready|43665       |61               |5/31/2011|6/12/2011|6/7/2011|283      |517       |14352.7713|1375.9427|429.9821|16158.6961|711      |2       |20.1865  |0.0000          |40.3730  |
|517       |Richard  |Bready  |Richard Bready|43665       |62               |5/31/2011|6/12/2011|6/7/2011|283      |517       |14352.7713|1375.9427|429.9821|16158.6961|773      |1       |2039.9940|0.0000          |2039.9940|
|517       |Richard  |Bready  |Richard Bready|43665       |63               |5/31/2011|6/12/2011|6/7/2011|283      |517       |14352.7713|1375.9427|429.9821|16158.6961|707      |1       |20.1865  |0.0000          |20.1865  |
|517       |Richard  |Bready  |Richard Bready|43665       |64               |5/31/2011|6/12/2011|6/7/2011|283      |517       |14352.7713|1375.9427|429.9821|16158.6961|715      |2       |28.8404  |0.0000          |57.6808  |
|517       |Richard  |Bready  |Richard Bready|43665       |65               |5/31/2011|6/12/2011|6/7/2011|283      |517       |14352.7713|1375.9427|429.9821|16158.6961|777      |2       |2024.9940|0.0000          |4049.9880|
|517       |Richard  |Bready  |Richard Bready|43665       |66               |5/31/2011|6/12/2011|6/7/2011|283      |517       |14352.7713|1375.9427|429.9821|16158.6961|712      |2       |5.1865   |0.0000          |
```

[33]:
```
#Get customers that only have surname Adams
dfAdams = sparkDf.where("lastname =='Adams'")

# inner join on Adams DF and orders
dfAdams.join(dfOrders,dfAdams.customerid ==  dfOrders.customerid,"inner").show()
```

```
+----------+---------+--------+---------+-----------+-----------------+---------+---------+---------+----------+----------+---------+---------+---------+---------+--------+---------+----------------+---------+
|customerid|firstname|lastname| fullname|salesorderid|salesorderdetailid| orderdate|  duedate| shipdate|employeeid|customerid|  subtotal|   taxamt|  freight|  totaldue|productid|orderqty|unitprice|unitpricediscount|linetotal|
+----------+---------+--------+---------+-----------+-----------------+---------+---------+---------+----------+----------+---------+---------+---------+---------+--------+---------+----------------+---------+
|       307|      Jay|   Adams|Jay Adams|       48382|            23857|10/30/2012|11/11/2012|11/6/2012|       277|       307|  20.5200|  2.0246|  0.6327|  23.1773|      805|       1|  20.5200|          0.0000|  20.5200|
|       307|      Jay|   Adams|Jay Adams|       50734|            34874| 4/30/2013| 5/12/2013| 5/7/2013|       275|       307| 2232.8181| 220.3047|  68.8452| 2521.9680|      739|       3| 744.2727|          0.0000|2232.8181|
|       307|      Jay|   Adams|Jay Adams|       53561|            49267| 7/31/2013| 8/12/2013| 8/7/2013|       275|       307|48717.0900|4705.2935|1470.4042|54892.7877|      952|       6|  12.1440|          0.0000|  72.8640|
|       307|      Jay|   Adams|Jay Adams|       53561|            49268| 7/31/2013| 8/12/2013| 8/7/2013|       275|       307|48717.0900|4705.2935|1470.4042|54892.7877|      739|       1| 818.7000|          0.0000| 818.7000|
|       307|      Jay|   Adams|Jay Adams|       53561|            49269| 7/31/2013| 8/12/2013| 8/7/2013|       275|       307|48717.0900|4705.2935|1470.4042|54892.7877|      985|       4| 338.9940|
```

[34]:
```
#Left join on orders and adams df
dfOrders.join(dfAdams,dfAdams.customerid ==  dfOrders.customerid,"left").show(100)
```

```
+----------+-----------------+---------+---------+---------+----------+----------+----------+---------+---------+----------+---------+-----------+-----------------+---------+----------+---------+--------+
|salesorderid|salesorderdetailid|orderdate|  duedate|shipdate|employeeid|customerid|  subtotal|   taxamt|  freight|  totaldue|productid|orderqty|unitprice|unitpricediscount|linetotal|customerid|firstname|lastname|fullname|
+----------+-----------------+---------+---------+---------+----------+----------+----------+---------+---------+----------+---------+-----------+-----------------+---------+----------+---------+--------+--------+
|       43665|                61|5/31/2011|6/12/2011|6/7/2011|       283|       517|14352.7713|1375.9427| 429.9821|16158.6961|      711|          2|  20.1865|           0.0000|  40.3730|      NULL|     NULL|    NULL|    NULL|
|       43665|                62|5/31/2011|6/12/2011|6/7/2011|       283|       517|14352.7713|1375.9427| 429.9821|16158.6961|      773|          1|2039.9940|           0.0000|2039.9940|      NULL|     NULL|    NULL|    NULL|
|       43665|                63|5/31/2011|6/12/2011|6/7/2011|       283|       517|14352.7713|1375.9427| 429.9821|16158.6961|      707|          1|  20.1865|           0.0000|  20.1865|      NULL|     NULL|    NULL|    NULL|
```

[36]:
```
# write down the data in converted Dynamic Frame to S3 location.
glueContext.write_dynamic_frame.from_options(
                                frame = dyfCustomersConvert,
                                connection_type="s3",
                                connection_options = {"path": "s3://sai-kishore-pyspark/write_down_dyf_to_s3_2"},
                                format = "csv",
                                format_options={
                                    "separator": ","
                                    },
                                transformation_ctx = "datasink2")
```

```
<awsglue.dynamicframe.DynamicFrame object at 0x7f277b41f6d0>
```

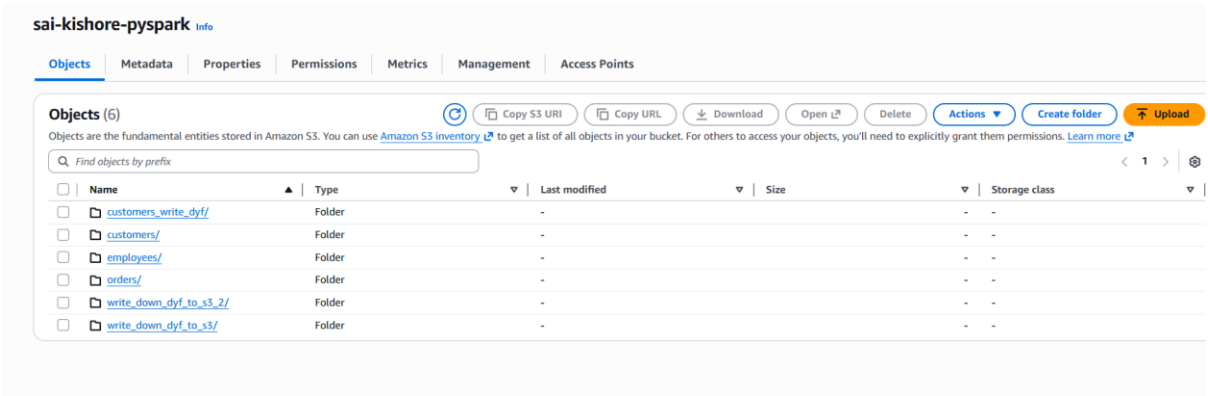

[38]:
```
# write data from the converted to customers_write_dyf table using the meta data stored in the glue data catalog
glueContext.write_dynamic_frame.from_catalog(
    frame = dyfCustomersConvert,
    database = "pyspark_tutorial_db",
    table_name = "customers_write_dyf")
```

```
<awsglue.dynamicframe.DynamicFrame object at 0x7f277b41cad0>
```

[ ]: